

# Introduction to Singularity

**Henric Zazzi**



# Overview

- What are containers
- Docker, the most popular container
- Singularity: Containers for the HPC environment
- installation of singularity
- How to build containers
- Installing software in container
- Running your container in an HPC environment
- Creating recipes for singularity


# What are containers






A container image is a lightweight, standalone, executable package of software that includes everything needed to run an application.

## Virtual Machine



## Container



# Containers: How are they useful

- Reproducibility
- Portability
- Depending on application and use-case, simple extreme scalability
- Next logical progression from virtual machines

# Why do we want containers in HPC?

- Escape “dependency hell”
- Load fewer modules
- Local and remote code works identically every time
- One file contains everything and can be moved anywhere

# Docker, the most popular container



# The Docker container software

- The most know and utilized container software
- Facilites workflow for creating, maintaining and distributing software
- Easy to install, well documented, standardized
- Used by many scientist



# Docker on HPC: The problem

- Incompabilities with scheduling managers (SLURM...)
- No support for MPI
- No native GPU support
- Docker users can escalate to root access on the cluster
- **Not allowed on HPC clusters**

# Singularity: Containers for the HPC environment

- Package software and dependencies in one file
- Use same container in different SNIC clusters
- Limits user's privileges, better security
- Same user inside container as on host
- No need for most modules
- Negligible performance decrease

## But I want to keep using docker

- Works great for local and private resources.
- No HPC centra will install docker for you
- Singularity can import Docker images



# Singularity hub

<https://singularity-hub.org/>



Container Collections

Want to build containers?

LOGIN

One collection is created for each connected Github repository. In that collection, several containers will automatically be built: one for each uniquely named recipe file found in the master branch of the Github repository.  
Read more about [recipe file naming](#) or [build options](#).

Enter Keywords Here

Name	Buids	Description	Stars ★	Downloads	Last Modified
<a href="#">researchapps/quantum_state_diffusion</a>	109	Solving quantum state diffusion numerically.	1	22	2017-10-19 🔗
<a href="#">vsoch/singularity-hello-world</a>	102		1	9098	2018-03-02 🔗
<a href="#">vsoch/pe-predictive</a>	100	run pe-finder to predict pumonary embolism from radiology reports	0	15	2017-10-17 🔗
<a href="#">TomHarrop/singularity-containers</a>	27		0	157	2018-09-10 🔗
<a href="#">dominik-handler/AP_singu</a>	16		0	621	2018-09-06 🔗
<a href="#">QuentinLetourneur/Let-it-bin</a>	15	Optimize workflow for binning metagenomic short reads from multiple samples	0	2	2018-09-10 🔗
<a href="#">marcc-hpc/tensorflow</a>	13	Singularity file wrapping Dockerfile of tensorflow/tensorflow:latest-gpu	1	4417	2018-09-09 🔗

# Singularity Versions

## **singularityCE (Community Edition)**

- Latest version: 3.11.0 (2023-02-17)
- Installed on Dardel: 3.10.4

# Singularity workflow

## Local computer

*Root access*

1. Create container
2. *singularity build*
3. Install software
4. Install libraries

## HPC Cluster

*User access*

1. *singularity shell*
2. *singularity exec*
3. *singularity help*
4. *singularity run*



# Install singularity on your computer

**You need a local installed copy of singularity to create your container**

```
$ sudo apt-get update && sudo apt-get install -y \  
    build-essential libssl-dev uuid-dev libgpgme11-dev \  
    squashfs-tools libseccomp-dev pkg-config \  
    libglib2.0-dev  
# Google GO language is needed  
$ sudo apt install golang  
$ git clone --recurse-submodules https://github.com/sylabs/singularity/  
$ cd singularity  
$ ./mconfig && make -C ./builddir && sudo make -C ./builddir install
```

More information at ...

<https://docs.sylabs.io/guides/3.0/user-guide/installation.html>

# Launching a container

- Singularity sets up the container environment and creates the necessary namespaces.
- Directories, files and other resources are shared from the host into the container.
- All expected I/O is passed through the container: pipes, program arguments, std, X11
- When the application(s) finish their foreground execution process, the container and namespaces collapse and vanish cleanly



# How to build containers

# Download and test an image

Download and test the latest UBUNTU image from docker hub

```
$ sudo singularity build my_image.sif docker://ubuntu:latest
INFO:      Starting build...
Getting image source signatures
...
INFO:      Creating SIF file...
INFO:      Build complete: my_image.sif
$ singularity shell my_image.sif
Singularity> cat /etc/*-release
DISTRIB_ID=Ubuntu
DISTRIB_RELEASE=22.04
DISTRIB_CODENAME=jammy
Singularity> exit
```

# Paths for building containers

```
singularity build [image].sif [name]
```

From	Path
Singularity hub	shub://[name]
Docker hub	docker://[name]
Local	[name]
Sandbox	[Sandbox folder name]
Recipe	[recipe name]

# How do I execute commands in singularity

Commands in the container can be given as normal.

```
$ singularity exec my_image.sif ls
```

```
$ singularity shell my_image.sif  
Singularity> ls
```

Do it yourself

## Exercise 1: Download a container

1. Go to singularity hub and find the hello-world container  
(<https://singularityhub.github.io/singularityhub-archive/>)
2. build the container using singularity
3. Use the container shell and get acquainted with it

# Installing software in container

# Why must I be root?

Same permissions in the container as outside...

To be root in the singularity image you must be root on the computer

# Build a writeable image

Since there are memory limitation on writing directly to image file, it is better to create a sandbox

```
$ sudo singularity build --sandbox my_sandbox my_image.sif
INFO:      Starting build...
INFO:      Verifying bootstrap image my_image.sif
...
INFO:      Creating sandbox directory...
INFO:      Build complete: my_sandbox
$ singularity shell my_sandbox
Singularity>
```



# Transfer files into container

**Read mode:** You can read/write to file system outside container and read inside container.

**write mode:** You can read/write inside container.

**Remember:** In write mode you are user ROOT, home folder: /root

# Binding folders

```
... -B [local folder]:[singularity folder] ...
```

- Enables transferring of data to container
- Enables accessing external data from within the container
- Enables running external scripts from within the container
- For using PDC filesystem you must bind to *cfs/klemming*

# How to use binding to run local scripts

## 1. Create local folder and internal singularity folder

```
$ mkdir my_folder  
$ sudo singularity exec -w my_sandbox mkdir /usr/local/sing
```

## 2. Starting container and bind folders

The file *myscript*, residing in *my\_folder*, will be executed as within the container

```
$ singularity exec -B my_folder:/usr/local/sing -w my_sandbox /usr/local/sing/myscript
```

## 3. Obscuring container folder */opt*

```
$ singularity exec -B my_folder:/opt my_sandbox /opt/myscript
```



# Example on how to transfer files into the container

1. Create local folder and internal singularity folder

```
$ mkdir my_folder  
$ sudo singularity exec -w my_sandbox mkdir singularity_folder
```

2. Starting container as *root* and bind folders

```
$ sudo singularity shell -B my_folder:/root/singularity_folder -w my_sandbox
```

3. Copy *file1* to container folder

```
Singularity> cp singularity_folder/file1 .
```

Do it yourself:

## Exercise 2: Create your own container

1. Go to docker hub and find the official latest ubuntu
2. build the container using singularity
3. Build a writeable sandbox
4. Install necessary tools into the container (Compiler etc...)
  - i. apt-get update
  - ii. apt-get install build-essential

# Initiating your container

# singularity.d folder

Startup scripts etc... for your singularity image

```
$ singularity exec my_image.sif ls -l /.singularity.d
-rw-r--r-- 1 root root    39 Feb 17 09:27 Singularity
drwxr-xr-x 2 root root 4096 Feb 17 09:27 actions
drwxr-xr-x 2 root root 4096 Feb 17 09:27 env
-rw-r--r-- 1 root root   459 Feb 17 09:27 labels.json
drwxr-xr-x 2 root root 4096 Feb 17 09:27 libs
-rwxr-xr-x 1 root root 1933 Feb 17 09:27 runscript
-rwxr-xr-x 1 root root  10 Feb 17 09:27 runscript.help
-rwxr-xr-x 1 root root   10 Feb 17 09:27 startscript
```

**Important: The files must be executable and owned by root**

# Creating a script

## Example of runscript file

```
#!/bin/sh  
echo "Hello world!"
```

## Executing the runscript file

```
$ singularity run my_image.sif  
Hello world!
```



# What is a help file and how is it used

## Example of runscript.help file

```
This is a text file
```

## Print the information within

```
$ singularity run-help my_image.sif  
This is a text file
```

Do it yourself:

## Exercise 3: Edit your own container

1. Create a help file
2. Create/Edit the runscript printing *Hello world!*

**Tip:** You can use an editor in your VM or create it and then transfer the file

# Creating recipes for singularity

# Singularity Recipes

A Singularity Recipe is the driver of a custom build, and the starting point for designing any custom container. It includes specifics about installation software, environment variables, files to add, and container metadata

# How to build from a recipe

A recipe is a textfile explaining what should be put into the container

```
sudo singularity build [container].sif [recipe].def
```

Recipes for images that can be used on PDC clusters can be found at

<https://github.com/PDC-support/PDC-SoftwareStack/tree/master/other/singularity>

# Printing how a container was built

```
singularity inspect --deffile [container]
```

# Recipe format

```
# Header
Bootstrap: docker
From: ubuntu:latest
# Sections
%help
    Help me. I'm in the container.
%files
    mydata.txt /home
%post
    apt-get -y update
    apt-get install -y build-essential
%runscript
    echo "This is my runscript"
```

# Header

What image should we start with?

- *Bootstrap:*
  - shub
  - docker
  - localimage
- *From:*
  - The name of the container

```
# Header
Bootstrap: docker
From: ubuntu:latest
```



## Section: %help

Some information about your container.

Valuable to put information about what software and versions are available in the container

```
%help
```

```
This container is based on UBUNTU 22.04. GCC installed
```

## Section: %post

What softwares should be installed in my container.

```
%post  
apt-get -y update  
apt-get install -y build-essential
```

No interaction in the scripts

We do not need sudo in the container

## Section: %files

What local files should be copied into my container

```
%files  
  <local filename> <singularity path>
```

## Section: %runscript

What should be executed with the run command.

```
%runscript  
  <software executable> -<parameter1>
```

Do it yourself:

## Exercise 4: Create a recipe

1. Based on UBUNTU
2. Install compilers
3. Create a help text
4. Create a runscript
5. Run the recipe

**Tip:** You can use the editor in your VM and then transfer the file

# Running your container in an HPC environment

# Requirements

- OpenMPI version must be the same in container and cluster
- Compiler and version must be the same in container and cluster
- You need to bind to the LUSTRE file system at PDC so it can be detected
- You can use *build* but only from other images and only sandboxes
- You can **ONLY** run *sandbox* and not *SIF* files
  - A singularity file is copied to all nodes whereas a *sandbox* folder structure is not

## Transfer of sandbox file

```
# On local computer
sudo tar czf <sandbox name>.tar.gz <sandbox name>
scp <sandbox name>.tar.gz <username>@dardel.pdc.kth.se:/cfs/klemming/home/<u>/<username>
# On Dardel
tar xfp <sandbox name>.tar.gz
```

## Transfer a SIF file

```
scp <SIF file> <username>@dardel.pdc.kth.se:/cfs/klemming/home/<u>/<username>
singularity build --sandbox <sandbox name> <SIF file>
```



# What are the required tools

- **Packages:** wget git bash gcc gfortran g++ make
- **Source:** MPICH

```
ml PDC  
ml singularity
```

- In folder `$PDC_SHUB` you can find already built images at PDC

# Executes singularity on 2 nodes

```
#!/bin/bash -l
# The -l above is required to get the full environment with modules
# Set the allocation to be charged for this job
#SBATCH -A 202X-X-XX
# The name of the script is myjob
#SBATCH -J myjob
# Only 1 hour wall-clock time will be given to this job
#SBATCH -t 1:00:00
# Number of nodes
#SBATCH --nodes=2
# Using the shared partition as we are not using all cores
#SBATCH -p shared
# Number of MPI processes per node
#SBATCH --ntasks-per-node=12
# Run the executable named myexe
ml PDC singularity
srun -n 24 --mpi=pmi2 singularity exec <sandbox folder> <myexe>
```



# Executes GPU enabled code with singularity

```
#!/bin/bash -l
# The -l above is required to get the full environment with modules
# Set the allocation to be charged for this job
# not required if you have set a default allocation
#SBATCH -A 201X-X-XX
# The name of the script is myjob
#SBATCH -J myjob
# Only 1 hour wall-clock time will be given to this job
#SBATCH -t 1:00:00
# Number of nodes
#SBATCH --nodes=1
# Using the GPU partition which is at the moment is under testing
#SBATCH -p gpu-tst
# Run the executable named myexe
ml PDC singularity
srun -n 1 singularity exec --rocm -B /cfs/klemming <sandbox folder> <myexe>
```

## Exercise 5: Run a HPC container

1. Login into `dardel.pdc.kth.se`
2. send in a job for the hello-world sandbox
3. Use the `hello_world` in PDCs singularity repository

**Tip:** With the singularity module use the **Path:** `$PDC_SHUB`

## Useful links

- <https://www.pdc.kth.se/support/documents/software/singularity.html>
- <https://github.com/PDC-support/PDC-SoftwareStack/tree/master/other/singularity>
- <https://sylabs.io/docs/>