

Introduction to Singularity

Henric Zazzi



Overview

- What are containers
- Docker, the most popular container
- Singularity: Containers for the HPC environment
- installation of singularity
- Using the container
- How to build containers
- Running your container in an HPC environment
- Creating recipes for singularity


What are containers






A container image is a lightweight, standalone, executable package of software that includes everything needed to run an application.

Virtual Machine



Container



Containers: How are they useful

- Reproducibility
- Portability
- Depending on application and use-case, simple extreme scalability
- Next logical progression from virtual machines

Why do we want containers in HPC?

- Escape “dependency hell”
- Load fewer modules
- Local and remote code works identically every time
- One file contains everything and can be moved anywhere

Docker, the most popular container



The Docker container software

- The most know and utilized container software
- Facilites workflow for creating, maintaining and distributing software
- Easy to install, well documented, standardized
- Used by many scientist

Docker on HPC: The problem

- Incompabilities with scheduling managers (SLURM...)
- No support for MPI
- No native GPU support
- Docker users can escalate to root access on the cluster
- **Not allowed on HPC clusters**

Singularity: Containers for the HPC environment

- Package software and dependencies in one file
- Use same container in different SNIC clusters
- Limits user's privileges, better security
- Same user inside container as on host
- No need for most modules
- Negligible performance decrease

But I want to keep using docker

- Works great for local and private resources.
- No HPC centra will install docker for you
- Singularity can import Docker images



Singularity hub

<https://singularity-hub.org/>



Container Collections

Want to build containers?

LOGIN

One collection is created for each connected Github repository. In that collection, several containers will automatically be built: one for each uniquely named recipe file found in the master branch of the Github repository.
Read more about [recipe file naming](#) or [build options](#).

Enter Keywords Here

Name	Buids	Description	Stars ★	Downloads	Last Modified
researchchapps/quantum_state_diffusion	109	Solving quantum state diffusion numerically.	1	22	2017-10-19 🔗
vsoch/singularity-hello-world	102		1	9098	2018-03-02 🔗
vsoch/pe-predictive	100	run pe-finder to predict pumonary embolism from radiology reports	0	15	2017-10-17 🔗
TomHarrop/singularity-containers	27		0	157	2018-09-10 🔗
dominik-handler/AP_singu	16		0	621	2018-09-06 🔗
QuentinLetourneur/Let-it-bin	15	Optimize workflow for binning metagenomic short reads from multiple samples	0	2	2018-09-10 🔗
marcc-hpc/tensorflow	13	Singularity file wrapping Dockerfile of tensorflow/tensorflow:latest-gpu	1	4417	2018-09-09 🔗

Singularity Versions

singularityCE (Community Edition)

- Latest version: 3.10.2 (2022-07-22)
- Installed on Dardel: 3.10.0

Singularity workflow

Local computer

Root access

1. Create container
2. *singularity build*
3. Install software
4. Install libraries

HPC Cluster

User access

1. *singularity shell*
2. *singularity exec*
3. *singularity help*
4. *singularity run*

Install singularity on your computer

You need a local installed copy of singularity to create your container

```
$ sudo apt-get update && sudo apt-get install -y \  
    build-essential libssl-dev uuid-dev libgpgme11-dev \  
    squashfs-tools libseccomp-dev pkg-config \  
    libglib2.0-dev  
# Google G0 language is needed  
$ sudo apt install golang  
$ git clone --recurse-submodules https://github.com/sylabs/singularity/  
$ cd singularity  
$ ./mconfig && make -C ./builddir && sudo make -C ./builddir install
```

More information at ...

<https://docs.sylabs.io/guides/3.0/user-guide/installation.html>

Launching a container

- Singularity sets up the container environment and creates the necessary namespaces.
- Directories, files and other resources are shared from the host into the container.
- All expected I/O is passed through the container: pipes, program arguments, std, X11
- When the application(s) finish their foreground execution process, the container and namespaces collapse and vanish cleanly

Using the container

Download and test an image

Download and test the latest UBUNTU image from docker hub

```
$ sudo singularity build my_image.sif docker://ubuntu:latest
INFO:      Starting build...
Getting image source signatures
...
INFO:      Creating SIF file...
INFO:      Build complete: my_image.sif
$ singularity shell my_image.sif
Singularity> cat /etc/*-release
DISTRIB_ID=Ubuntu
DISTRIB_RELEASE=22.04
DISTRIB_CODENAME=jammy
Singularity> exit
```

Do it yourself

Exercise 1: Download a container

1. Go to singularity hub and find the hello-world container (<https://singularity-hub.org/collections>)
2. build the container using singularity
(shub://[full name of container])
3. Use the container shell and get acquainted with it

How to build containers

Why must I be root?

Same permissions in the container as outside...

To be root in the singularity image you must be root on the computer

Build a writeable image

Since there are memory limitation on writing directly to image file, it is better to create a sandbox

```
$ sudo singularity build --sandbox my_sandbox my_image.sif
INFO:      Starting build...
INFO:      Verifying bootstrap image my_image.sif
...
INFO:      Creating sandbox directory...
INFO:      Build complete: my_sandbox
$ singularity shell my_sandbox
Singularity>
```

How do I execute commands in singularity

Commands in the container can be given as normal.

```
$ singularity exec my_image.sif ls
```

```
$ singularity shell my_image.sif  
Singularity> ls
```

Transfer files into container

Read mode: You can read/write to file system outside container and read inside container.

write mode: You can read/write inside container.

Remember: In write mode you are user ROOT, home folder: /root

How to transfer files into the container

```
$ sudo singularity exec -w my_sandbox mkdir singularity_folder  
$ mkdir my_folder  
$ sudo singularity shell -B my_folder:/root/singularity_folder -w my_sandbox  
Singularity> cp singularity_folder/file1 .
```

Do it yourself:

Exercise 2: Create your own container

1. Go to docker hub and find the official latest ubuntu
2. build the container using singularity
3. Build a writeable sandbox
4. Install necessary tools into the container (Compiler etc...)
 - i. apt-get update
 - ii. apt-get install build-essential

singularity.d folder

Startup scripts etc... for your singularity image

```
$ singularity exec my_image.sif ls -l /.singularity.d
total 1
drwxr-xr-x 2 root root 76 Sep 11 17:05 actions
drwxr-xr-x 2 root root 139 Sep 11 17:23 env
drwxr-xr-x 2 root root 3 Sep 11 17:05 libs
-rwxr-xr-x 1 root root 33 Sep 11 17:23 runscript
-rwxr-xr-x 1 root root 10 Sep 11 17:05 runscript.help
```

Important: The files must be executable and owned by root

Creating a script

runscript

```
#!/bin/sh  
  
ls -l
```

command

```
$ singularity run my_image.sif  
total 1  
drwxr-xr-x 2 root root 76 Sep 11 17:05 file1  
drwxr-xr-x 2 root root 139 Sep 11 17:23 file2
```

What is a help file and how is it used

runscript.help

```
This is a text file
```

command

```
$ singularity run-help my_image.sif  
This is a text file
```

Build a new container from a sandbox

```
$ sudo singularity build my_new_image.sif my_sandbox
INFO:      Starting build...
INFO:      Creating SIF file...
INFO:      Build complete: my_new_image.sif
```

Do it yourself:

Exercise 3: Edit your own container

1. Create a help file
2. Create/Edit the runscript running hello world
3. Create a new container from the sandbox

Tip: You can use the editor in your VM and then transfer the file

Creating recipes for singularity

Singularity Recipes

A Singularity Recipe is the driver of a custom build, and the starting point for designing any custom container. It includes specifics about installation software, environment variables, files to add, and container metadata

How to build from a recipe

A recipe is a textfile explaining what should be put into the container

```
sudo singularity build my_image.sif my_recipe
```

Recipes for images that can be used on PDC clusters can be found at

<https://github.com/PDC-support/PDC-SoftwareStack/tree/master/other/singularity>

Recipe format

```
# Header
Bootstrap: docker
From: ubuntu:latest
# Sections
%help
    Help me. I'm in the container.
%files
    mydata.txt /home
%post
    apt-get -y update
    apt-get install -y build-essential
%runscript
    echo "This is my runscript"
```

Header

What image should we start with?

- *Bootstrap:*
 - shub
 - docker
 - localimage
- *From:*
 - The name of the container

```
# Header
Bootstrap: docker
From: ubuntu:latest
```

Section: %help

Some information about your container.

Valuable to put information about what software and versions are available in the container

```
%help
```

```
This container is based on UBUNTU 22.04. GCC installed
```

Section: %post

What softwares should be installed in my container.

```
%post
apt-get -y update
apt-get install -y build-essential
```

No interaction in the scripts

We do not need sudo in the container

Section: %files

What local files should be copied into my container

```
%files
# <filename> <singularity path>
myfile.txt /opt
```

Section: %runscript

What should be executed with the run command.

```
%runscript  
mysoftware -param1 -param2
```

Do it yourself:

Exercise 5: Create a recipe

1. Based on UBUNTU
2. Install compilers
3. Create a help text
4. Create a runscript
5. Run the recipe

Tip: You can use the editor in your VM and then transfer the file

Running your container in an HPC environment

Requirements

- OpenMPI version must be the same in container and cluster
- Compiler and version must be the same in container and cluster
- You need to bind to the LUSTRE file system at PDC so it can be detected
- You can use *build* but only from other images and only sandboxes
- You can **ONLY** run *sandbox* and not *SIF* files
 - A singularity file is copied to all nodes whereas a *sandbox* folder structure is not

Transfer of sandbox file

```
# On local computer
sudo tar czf <sandbox name>.tar.gz <sandbox name>
scp <sandbox name>.tar.gz <username>@dardel.pdc.kth.se:/cfs/klemming/home/<u>/<username>
# On Dardel
tar xfp <sandbox name>.tar.gz
```

What are the required tools

- **Packages:** wget git bash gcc gfortran g++ make
- **Source:** MPICH

```
ml PDC  
ml singularity
```

- In folder `$PDC_SHUB` you can find already built images at PDC
- <https://www.pdc.kth.se/software>

Executes singularity on 2 nodes

```
#!/bin/bash -l
# The -l above is required to get the full environment with modules
# Set the allocation to be charged for this job
#SBATCH -A 202X-X-XX
# The name of the script is myjob
#SBATCH -J myjob
# Only 1 hour wall-clock time will be given to this job
#SBATCH -t 1:00:00
# Number of nodes
#SBATCH --nodes=2
# Using the shared partition as we are not using all cores
#SBATCH -p shared
# Number of MPI processes per node
#SBATCH --ntasks-per-node=12
# Run the executable named myexe
ml PDC singularity
srun -n 24 --mpi=pmi2 singularity exec <sandbox folder> <myexe>
```

Do it yourself:

Exercise 4: Run a HPC container

1. Login into dardel.pdc.kth.se
2. send in a job for the hello-world sandbox
3. Use the hello_world in PDCs singularity repository

Tip: With the singularity module use the **Path:** \$PDC_SHUB

Useful links

- <https://www.pdc.kth.se/software/software/singularity/>
- <https://docs.sylabs.io/guides/3.8/user-guide/>
- <https://github.com/PDC-support/PDC-SoftwareStack/tree/master/other/singularity>